

An RNN model for temperature data

This time we will be working with real data: daily (Tmin, Tmax) temperature series from 36 weather stations spanning 50 years. It is to be noted that a pretty good predictor model already exists for temperatures: the average of temperatures on the same day of the year in N previous years. It is not clear if RNNs can do better but we will see how far they can go.

In [1]:

```
!pip install tensorflow==1.15.3
```

```
Collecting tensorflow==1.15.3
  Downloading tensorflow-1.15.3-cp37-cp37m-manylinux2010_x86_64.whl (110.5 MB)
    |████████████████████████████████████████| 110.5 MB 22 kB/s s eta 0:00:01
Collecting gast==0.2.2
  Downloading gast-0.2.2.tar.gz (10 kB)
Collecting tensorflow-estimator==1.15.1
  Downloading tensorflow_estimator-1.15.1-py2.py3-none-any.whl (503 kB)
    |████████████████████████████████████████| 503 kB 49.8 MB/s eta 0:00:01
Requirement already satisfied: protobuf>=3.6.1 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (3.18.1)
Requirement already satisfied: absl-py>=0.7.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (0.14.1)
Collecting keras-applications>=1.0.8
  Downloading Keras_Applications-1.0.8-py3-none-any.whl (50 kB)
    |████████████████████████████████████████| 50 kB 8.8 MB/s eta 0:00:01
Requirement already satisfied: wrapt>=1.11.1 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (1.13.1)
Requirement already satisfied: grpcio>=1.8.6 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (1.38.1)
Requirement already satisfied: numpy<2.0,>=1.16.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (1.19.5)
Requirement already satisfied: opt-einsum>=2.3.2 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (3.3.0)
Requirement already satisfied: six>=1.10.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (1.16.0)
Requirement already satisfied: google-pasta>=0.1.6 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (0.2.0)
Collecting tensorboard<1.16.0,>=1.15.0
  Downloading tensorboard-1.15.0-py3-none-any.whl (3.8 MB)
    |████████████████████████████████████████| 3.8 MB 16.4 MB/s eta 0:00:01
Requirement already satisfied: wheel>=0.26 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (0.37.0)
Collecting astor>=0.6.0
  Downloading astor-0.8.1-py2.py3-none-any.whl (27 kB)
Requirement already satisfied: termcolor>=1.1.0 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (1.1.0)
Requirement already satisfied: keras-preprocessing>=1.0.5 in /opt/conda/lib/python3.7/site-packages (from tensorflow==1.15.3) (1.1.2)
Requirement already satisfied: h5py in /opt/conda/lib/python3.7/site-packages (from keras-applications>=1.0.8->tensorflow==1.15.3) (3.1.0)
Requirement already satisfied: markdown>=2.6.8 in /opt/conda/lib/python3.7/site-packages (from tensorboard<1.16.0,>=1.15.0->tensorflow==1.15.3) (3.3.4)
Requirement already satisfied: werkzeug>=0.11.15 in /opt/conda/lib/python3.7/site-packages (from tensorboard<1.16.0,>=1.15.0->tensorflow==1.15.3) (2.0.2)
Requirement already satisfied: setuptools>=41.0.0 in /opt/conda/lib/python3.7/site-packages (from tensorboard<1.16.0,>=1.15.0->tensorflow==1.15.3) (58.2.0)
Requirement already satisfied: importlib-metadata in /opt/conda/lib/python3.7/si
```

```

te-packages (from markdown>=2.6.8->tensorboard<1.16.0,>=1.15.0->tensorflow==1.15.3) (4.8.1)
Requirement already satisfied: cached-property in /opt/conda/lib/python3.7/site-packages (from h5py->keras-applications>=1.0.8->tensorflow==1.15.3) (1.5.2)
Requirement already satisfied: typing-extensions>=3.6.4 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata->markdown>=2.6.8->tensorboard<1.16.0,>=1.15.0->tensorflow==1.15.3) (3.10.0.2)
Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata->markdown>=2.6.8->tensorboard<1.16.0,>=1.15.0->tensorflow==1.15.3) (3.6.0)
Building wheels for collected packages: gast
  Building wheel for gast (setup.py) ... done
  Created wheel for gast: filename=gast-0.2.2-py3-none-any.whl size=7554 sha256=08cb902cd0c6bdf6f5fc7731880608f137a80b3322f81bdc49678468b4b8b8b
  Stored in directory: /home/jupyter/.cache/pip/wheels/21/7f/02/420f32a803f7d0967b48dd823da3f558c5166991bfd204eef3
Successfully built gast
Installing collected packages: tensorflow-estimator, tensorboard, keras-applications, gast, astor, tensorflow
  Attempting uninstall: tensorflow-estimator
    Found existing installation: tensorflow-estimator 2.6.0
    Uninstalling tensorflow-estimator-2.6.0:
      Successfully uninstalled tensorflow-estimator-2.6.0
  Attempting uninstall: tensorboard
    Found existing installation: tensorboard 2.5.0
    Uninstalling tensorboard-2.5.0:
      Successfully uninstalled tensorboard-2.5.0
  Attempting uninstall: gast
    Found existing installation: gast 0.4.0
    Uninstalling gast-0.4.0:
      Successfully uninstalled gast-0.4.0
  Attempting uninstall: tensorflow
    Found existing installation: tensorflow 2.6.0
    Uninstalling tensorflow-2.6.0:
      Successfully uninstalled tensorflow-2.6.0
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
tensorflow-io 0.18.0 requires tensorflow-io-gcs-filesystem==0.18.0, which is not installed.
explainable-ai-sdk 1.3.2 requires xai-image-widget, which is not installed.
tfx-bsl 1.3.0 requires absl-py<0.13,>=0.9, but you have absl-py 0.14.1 which is incompatible.
tfx-bsl 1.3.0 requires google-api-python-client<2,>=1.7.11, but you have google-api-python-client 2.24.0 which is incompatible.
tfx-bsl 1.3.0 requires pyarrow<3,>=1, but you have pyarrow 5.0.0 which is incompatible.
tensorflow-transform 1.3.0 requires absl-py<0.13,>=0.9, but you have absl-py 0.14.1 which is incompatible.
tensorflow-transform 1.3.0 requires pyarrow<3,>=1, but you have pyarrow 5.0.0 which is incompatible.
tensorflow-serving-api 2.6.0 requires tensorflow<3,>=2.6.0, but you have tensorflow 1.15.3 which is incompatible.
tensorflow-probability 0.13.0rc0 requires gast>=0.3.2, but you have gast 0.2.2 which is incompatible.
tensorflow-io 0.18.0 requires tensorflow<2.6.0,>=2.5.0, but you have tensorflow 1.15.3 which is incompatible.
tensorflow-cloud 0.1.14 requires tensorboard>=2.3.0, but you have tensorboard 1.15.0 which is incompatible.
Successfully installed astor-0.8.1 gast-0.2.2 keras-applications-1.0.8 tensorboard-1.15.0 tensorflow-1.15.3 tensorflow-estimator-1.15.1

```

Please ignore any compatibility warnings and errors. Make sure to **restart** your kernel to ensure this change has taken place.

In [1]:

```
import math
import sys
import time
import numpy as np

sys.path.insert(0, 'temperatures/utils/') #so python can find the utils_ modules
import utils_batching
import utils_args

import tensorflow as tf
from tensorflow.python.lib.io import file_io as gfile
print("Tensorflow version: " + tf.__version__)
```

Tensorflow version: 1.15.3

In [2]:

```
from matplotlib import pyplot as plt
import utils_prettystyle
import utils_display
```

Download Data

In [3]:

```
%%bash
DOWNLOAD_DIR=temperatures/data
mkdir $DOWNLOAD_DIR
gsutil -m cp gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequ
```

```
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00025412.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00054452.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00033862.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00049742.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00026132.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00055722.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00013870.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014860.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014850.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014820.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00091732.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00013880.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USC00087982.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
```

```

temperatures/USW00003820.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00013960.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00013970.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014840.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00013740.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014740.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014750.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014940.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00014950.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00023050.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00023070.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00023160.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00023170.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00024090.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00024130.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00024150.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00026510.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00026410.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00093720.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00093820.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00094830.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00094910.csv...
Copying gs://cloud-training-demos/courses/machine_learning/deepdive/09_sequence/
temperatures/USW00093730.csv...
- [36/36 files][ 14.2 MiB/ 14.2 MiB] 100% Done
Operation completed over 36 objects/14.2 MiB.

```

Hyperparameters

N_FORWARD = 1: works but model struggles to predict from some positions

N_FORWARD = 4: better but still bad occasionnally

N_FORWARD = 8: works perfectly

```

In [4]: NB_EPOCHS = 5          # number of times the model sees all the data during trainin
        N_FORWARD = 8          # train the network to predict N in advance (traditionnally
        RESAMPLE_BY = 5        # averaging period in days (training on daily data is too mu

```

```

RNN_CELL_SIZE = 128 # size of the RNN cells
N_LAYERS = 2 # number of stacked RNN cells (needed for tensor shapes but
SEQUEN = 128 # unrolled sequence length
BATCH_SIZE = 64 # mini-batch size
DROPOUT_PKEEP = 0.7 # probability of neurons not being dropped (should be between
ACTIVATION = tf.nn.tanh # Activation function for GRU cells (tf.nn.relu or tf.nn

JOB_DIR = "checkpoints"
DATA_DIR = "temperatures/data"

# potentially override some settings from command-line arguments
if __name__ == '__main__':
    JOB_DIR, DATA_DIR = utils_args.read_args1(JOB_DIR, DATA_DIR)

ALL_FILEPATTERN = DATA_DIR + "/*.csv" # pattern matches all 1666 files
EVAL_FILEPATTERN = DATA_DIR + "/USC000*2.csv" # pattern matches 8 files
# pattern USW*.csv -> 298 files, pattern USW*0.csv -> 28 files
print('Reading data from "{}".\nWriting checkpoints to "{}".'.format(DATA_DIR,

```

Reading data from "temperatures/data".
Writing checkpoints to "checkpoints".

Temperature data

This is what our temperature datasets looks like: sequences of daily (Tmin, Tmax) from 1960 to 2010. They have been cleaned up and eventual missing values have been filled by interpolation. Interpolated regions of the dataset are marked in red on the graph.

In [5]:

```

all_filenames = gfile.get_matching_files(ALL_FILEPATTERN)
eval_filenames = gfile.get_matching_files(EVAL_FILEPATTERN)
train_filenames = list(set(all_filenames) - set(eval_filenames))

# By default, this utility function loads all the files and places data
# from them as-is in an array, one file per line. Later, we will use it
# to shape the dataset as needed for training.
ite = utils_batching.rnn_multistation_sampling_temperature_sequencer(eval_filenames,
evtemps, _, evdates, _, _ = next(ite) # gets everything

print('Pattern "{}" matches {} files'.format(ALL_FILEPATTERN, len(all_filenames)))
print('Pattern "{}" matches {} files'.format(EVAL_FILEPATTERN, len(eval_filenames)))
print("Evaluation files: {}".format(len(eval_filenames)))
print("Training files: {}".format(len(train_filenames)))
print("Initial shape of the evaluation dataset: " + str(evtemps.shape))
print("{} files, {} data points per file, {} values per data point"
      " (Tmin, Tmax, is_interpolated)".format(evtemps.shape[0], evtemps.shape[1]

```

Loading 8 files.....
Pattern "temperatures/data/*.csv" matches 36 files
Pattern "temperatures/data/USC000*2.csv" matches 8 files
Evaluation files: 8
Training files: 28
Initial shape of the evaluation dataset: (8, 18262, 3)
8 files, 18262 data points per file, 3 values per data point (Tmin, Tmax, is_interpolated)

In [6]:

```


# You can adjust the visualisation range and dataset here.
# Interpolated regions of the dataset are marked in red.

```



```
WEATHER_STATION = 0 # 0 to 7 in default eval dataset
START_DATE = 0      # 0 = Jan 2nd 1950
END_DATE = 18262     # 18262 = Dec 31st 2009
visu_temperatures = evtemps[WEATHER_STATION, START_DATE:END_DATE]
visu_dates = evdates[START_DATE:END_DATE]

utils_display.picture_this_4(visu_temperatures, visu_dates)
```



Resampling

Our RNN would need to be unrolled across 365 steps to capture the yearly temperature cycles. That's a bit too much. We will resample the temperatures and work with 5-day averages for example. This is what resampled (Tmin, Tmax) temperatures look like.


In [7]:

```
# This time we ask the utility function to average temperatures over 5-day period
ite = utils_batching.rnn_multistation_sampling_temperature_sequencer(eval_filenames, evaltemps, _, evaldates, _, _ = next(ite))
```

Loading 8 files.....

In [8]:

```
# display five years worth of data
WEATHER_STATION = 0      # 0 to 7 in default eval dataset
START_DATE = 0           # 0 = Jan 2nd 1950
END_DATE = 365*5//RESAMPLE_BY # 5 years
visu_temperatures = evaltemps[WEATHER_STATION, START_DATE:END_DATE]
visu_dates = evaldates[START_DATE:END_DATE]
plt.fill_between(visu_dates, visu_temperatures[:,0], visu_temperatures[:,1])
plt.show()
```



Visualize training sequences

This is what the neural network will see during training.

In [9]:

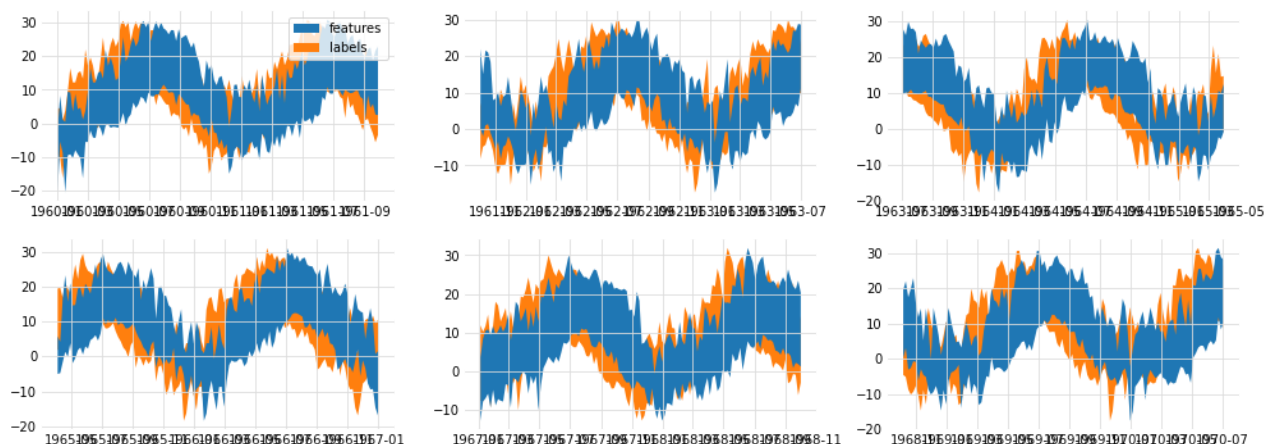
```
# The function rnn_multistation_sampling_temperature_sequencer puts one weather
# a batch and continues with data from the same station in corresponding lines i
# Features and labels are returned with shapes [BATCHSIZE, SEQLen, 2]. The last
# contains (Tmin, Tmax).
ite = utils_batching.rnn_multistation_sampling_temperature_sequencer(eval_filena
                                                                    RESAMPLE_BY
                                                                    BATCHSIZE,
                                                                    SEQLen,
                                                                    N_FORWARD,
                                                                    nb_epochs=1
                                                                    tminmax=True

# load 6 training sequences (each one contains data for all weather stations)
visu_data = [next(ite) for _ in range(6)]
```

Loading 8 files.....

In [10]:

```
# Check that consecutive training sequences from the same weather station are in
WEATHER_STATION = 4
utils_display.picture_this_5(visu_data, WEATHER_STATION)
```



The model definition



deep RNN schematic

```

X shape [BATCHSIZE, SEQLen, 2]
Y shape [BATCHSIZE, SEQLen, 2]
H shape [BATCHSIZE, RNN_CELLsize*NLAYERS]

```

When executed, this function instantiates the Tensorflow graph for our model.

In [11]:

```

def model_rnn_fn(features, Hin, labels, step, dropout_pkeep):
    X = features # shape [BATCHSIZE, SEQLen, 2], 2 for (Tmin, Tmax)
    batchsize = tf.shape(X)[0]
    seqlen = tf.shape(X)[1]
    pairlen = tf.shape(X)[2] # should be 2 (tmin, tmax)

    cells = [tf.nn.rnn_cell.GRUCell(RNN_CELLsize, activation=ACTIVATION) for _ in range(NLAYERS)]
    # dropout useful between cell layers only: no output dropout on last cell
    cells = [tf.nn.rnn_cell.DropoutWrapper(cell, output_keep_prob = dropout_pkeep) for cell in cells]
    # a stacked RNN cell still works like an RNN cell
    cell = tf.nn.rnn_cell.MultiRNNCell(cells, state_is_tuple=False)
    # X[BATCHSIZE, SEQLen, 2], Hin[BATCHSIZE, RNN_CELLsize*NLAYERS]
    # the sequence unrolling happens here
    Yn, H = tf.nn.dynamic_rnn(cell, X, initial_state=Hin, dtype=tf.float32)
    # Yn[BATCHSIZE, SEQLen, RNN_CELLsize]
    Yn = tf.reshape(Yn, [batchsize*seqlen, RNN_CELLsize])
    Yr = tf.layers.dense(Yn, 2) # Yr [BATCHSIZE*SEQLen, 2]
    Yr = tf.reshape(Yr, [batchsize, seqlen, 2]) # Yr [BATCHSIZE, SEQLen, 2]
    Yout = Yr[:, -N_FORWARD :, :] # Last N_FORWARD outputs Yout [BATCHSIZE, N_FORWARD, 2]

    loss = tf.losses.mean_squared_error(Yr, labels) # labels[BATCHSIZE, SEQLen, 2]

    lr = 0.001 + tf.train.exponential_decay(0.01, step, 1000, 0.5)
    optimizer = tf.train.AdamOptimizer(learning_rate=lr)
    train_op = optimizer.minimize(loss)

    return Yout, H, loss, train_op, Yr

```

Instantiate the model

In [12]:

```

tf.reset_default_graph() # restart model graph from scratch

# placeholder for inputs
Hin = tf.placeholder(tf.float32, [None, RNN_CELLsize * NLAYERS])
features = tf.placeholder(tf.float32, [None, None, 2]) # [BATCHSIZE, SEQLen, 2]
labels = tf.placeholder(tf.float32, [None, None, 2]) # [BATCHSIZE, SEQLen, 2]
step = tf.placeholder(tf.int32)
dropout_pkeep = tf.placeholder(tf.float32)

# instantiate the model
Yout, H, loss, train_op, Yr = model_rnn_fn(features, Hin, labels, step, dropout_pkeep)

```

WARNING:tensorflow:From /tmp/ipykernel_25662/71392446.py:1: The name tf.reset_default_graph is deprecated. Please use tf.compat.v1.reset_default_graph instead.

WARNING:tensorflow:From /tmp/ipykernel_25662/71392446.py:4: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:7: GRUCell.__init__

(from tensorflow.python.ops.rnn_cell_impl) is deprecated and will be removed in a future version.

Instructions for updating:

This class is equivalent as tf.keras.layers.GRUCell, and will be replaced by that in Tensorflow 2.0.

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:9: The name tf.nn.rnn_cell.DropoutWrapper is deprecated. Please use tf.compat.v1.nn.rnn_cell.DropoutWrapper instead.

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:11: MultiRNNCell.__init__ (from tensorflow.python.ops.rnn_cell_impl) is deprecated and will be removed in a future version.

Instructions for updating:

This class is equivalent as tf.keras.layers.StackedRNNCells, and will be replaced by that in Tensorflow 2.0.

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:14: dynamic_rnn (from tensorflow.python.ops.rnn) is deprecated and will be removed in a future version.

Instructions for updating:

Please use `keras.layers.RNN(cell)`, which is equivalent to this API

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/rnn_cell_impl.py:559: Layer.add_variable (from tensorflow.python.keras.engine.base_layer) is deprecated and will be removed in a future version.

Instructions for updating:

Please use `layer.add_weight` method instead.

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/rnn_cell_impl.py:565: calling Constant.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/rnn_cell_impl.py:575: calling Zeros.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:17: dense (from tensorflow.python.layers.core) is deprecated and will be removed in a future version.

Instructions for updating:

Use keras.layers.Dense instead.

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/layers/core.py:187: Layer.apply (from tensorflow.python.keras.engine.base_layer) is deprecated and will be removed in a future version.

Instructions for updating:

Please use `layer.__call__` method instead.

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:21: The name tf.losses.mean_squared_error is deprecated. Please use tf.compat.v1.losses.mean_squared_error instead.

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/losses/losses_impl.py:121: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:23: The name tf.train.exponential_decay is deprecated. Please use tf.compat.v1.train.exponential_decay instead.

WARNING:tensorflow:From /tmp/ipykernel_25662/2393688187.py:24: The name tf.train

`n.AdamOptimizer` is deprecated. Please use `tf.compat.v1.train.AdamOptimizer` instead.

Initialize Tensorflow session

This resets all neuron weights and biases to initial random values

In [13]:

```
# variable initialization
sess = tf.Session()
init = tf.global_variables_initializer()
sess.run([init])
saver = tf.train.Saver(max_to_keep=1)
```

WARNING:tensorflow:From /tmp/ipykernel_25662/1210481336.py:2: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

WARNING:tensorflow:From /tmp/ipykernel_25662/1210481336.py:3: The name tf.global_variables_initializer is deprecated. Please use tf.compat.v1.global_variables_initializer instead.

WARNING:tensorflow:From /tmp/ipykernel_25662/1210481336.py:5: The name tf.train.Saver is deprecated. Please use tf.compat.v1.train.Saver instead.

2021-10-25 19:22:59.059977: I tensorflow/core/platform/cpu_feature_guard.cc:142] Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX2 FMA

2021-10-25 19:22:59.068433: I tensorflow/core/platform/profile_utils/cpu_utils.cc:94] CPU Frequency: 2199995000 Hz

2021-10-25 19:22:59.069358: I tensorflow/compiler/xla/service/service.cc:168] XLA service 0x562186374cd0 initialized for platform Host (this does not guarantee that XLA will be used). Devices:

2021-10-25 19:22:59.069405: I tensorflow/compiler/xla/service/service.cc:176] StreamExecutor device (0): Host, Default Version

The training loop

You can re-execute this cell to continue training.

Training data must be batched correctly, one weather station per line, continued on the same line across batches. This way, output states computed from one batch are the correct input states for the next batch. The provided utility function

`rnn_multistation_sampling_temperature_sequencer` does the right thing.

 batching for RNNs

In [14]:

```
losses = []
indices = []
last_epoch = 99999
last_fileid = 99999

for i, (next_features, next_labels, dates, epoch, fileid) in enumerate(
    utils_batching.rnn_multistation_sampling_temperature_sequencer(train_filenames,
                                                                    RESAMPLE_BY,
                                                                    BATCHSIZE,
```

```

SEQUEN,
N_FORWARD,
NB_EPOCHS, tm

# reinitalize state between epochs or when starting on data from a new wea
if epoch != last_epoch or fileid != last_fileid:
    batchsize = next_features.shape[0]
    H_ = np.zeros([batchsize, RNN_CELL_SIZE * N_LAYERS])
    print("State reset")

#train
feed = {Hin: H_, features: next_features, labels: next_labels, step: i, drop
Yout_, H_, loss_, _, Yr_ = sess.run([Yout, H, loss, train_op, Yr], feed_dict

# print progress
if i%20 == 0:
    print("{}: epoch {} loss = {} ({} weather stations this epoch)".format(i
sys.stdout.flush()
if i%10 == 0:
    losses.append(np.mean(loss_))
    indices.append(i)

last_epoch = epoch
last_fileid = fileid

# save the trained model
SAVEDMODEL = JOB_DIR + "/ckpt" + str(int(time.time()))
tf.saved_model.simple_save(sess, SAVEDMODEL,
                           inputs={"features":features, "Hin":Hin, "dropout_pkee
                           outputs={"Yout":Yout, "H":H})

```

Loading 28 files.....

State reset

0: epoch 0 loss = 294.3005676269531 (28 weather stations this epoch)

20: epoch 0 loss = 35.1406364440918 (28 weather stations this epoch)

State reset

40: epoch 1 loss = 26.80422019958496 (28 weather stations this epoch)

State reset

60: epoch 2 loss = 21.43130111694336 (28 weather stations this epoch)

80: epoch 2 loss = 23.1910400390625 (28 weather stations this epoch)

State reset

100: epoch 3 loss = 21.73704719543457 (28 weather stations this epoch)

State reset

120: epoch 4 loss = 21.566041946411133 (28 weather stations this epoch)

WARNING:tensorflow:From /tmp/ipykernel_25662/347898.py:39: simple_save (from tensorflow.python.saved_model.simple_save) is deprecated and will be removed in a future version.

Instructions for updating:

This function will only be available through the v1 compatibility library as tf.compat.v1.saved_model.simple_save.

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/saved_model/signature_def_utils_impl.py:201: build_tensor_info (from tensorflow.python.saved_model.utils_impl) is deprecated and will be removed in a future version.

Instructions for updating:

This function will only be available through the v1 compatibility library as tf.compat.v1.saved_model.utils.build_tensor_info or tf.compat.v1.saved_model.build_tensor_info.


INFO:tensorflow:Assets added to graph.

INFO:tensorflow:No assets to write.

INFO:tensorflow:SavedModel written to: checkpoints/ckpt1635189825/saved_model.pb

In [15]:

```
plt.ylim(ymax=np.amax(losses[1:])) # ignore first value for scaling
plt.plot(indices, losses)
plt.show()
```



Inference

This is a generative model: run an trained RNN cell in a loop

In [16]:

```
def prediction_run(predict_fn, prime_data, run_length):
    H = np.zeros([1, RNN_CELL_SIZE * N_LAYERS]) # zero state initially
    Yout = np.zeros([1, N_FORWARD, 2])
    data_len = prime_data.shape[0]-N_FORWARD

    # prime the state from data
    if data_len > 0:
        Yin = np.array(prime_data[:-N_FORWARD])
        Yin = np.reshape(Yin, [1, data_len, 2]) # reshape as one sequence of pairs
        r = predict_fn({'features': Yin, 'Hin':H, 'dropout_pkeep':1.0}) # no dropout
        Yout = r["Yout"]
        H = r["H"]

    # initailly, put real data on the inputs, not predictions
    Yout = np.expand_dims(prime_data[-N_FORWARD:], axis=0)
    # Yout shape [1, N_FORWARD, 2]: batch of a single sequence of length N_FORWARD

    # run prediction
    # To generate a sequence, run a trained cell in a loop passing as input and
    # respectively the output and output state from the previous iteration.
    results = []
    for i in range(run_length//N_FORWARD+1):
        r = predict_fn({'features': Yout, 'Hin':H, 'dropout_pkeep':1.0}) # no dropout
        Yout = r["Yout"]
        H = r["H"]
        results.append(Yout[0]) # shape [N_FORWARD, 2]

    return np.concatenate(results, axis=0)[:run_length]
```

Validation

In [17]:

```

QYEAR = 365// (RESAMPLE_BY*4)
YEAR = 365// (RESAMPLE_BY)

# Try starting predictions from January / March / July (resp. OFFSET = YEAR or Y
# Some start dates are more challenging for the model than others.
OFFSET = 30*YEAR+1*QYEAR

PRIMELEN=5*YEAR
RUNLEN=3*YEAR
RMSELEN=3*365// (RESAMPLE_BY*2) # accuracy of predictions 1.5 years in advance

```

In [18]:

```

# Restore the model from the last checkpoint saved previously.

# Alternative checkpoints:
# Once you have trained on all 1666 weather stations on Google Cloud ML Engine,
# SAVEDMODEL = "gs://{BUCKET}/sinejobs/sines_XXXXXX_XXXXXX/ckptXXXXXXXX"
# A sample checkpoint is provided with the lab. You can try loading it for compa
# SAVEDMODEL = "temperatures_best_checkpoint"

predict_fn = tf.contrib.predictor.from_saved_model(SAVEDMODEL)

```

WARNING:tensorflow:

The TensorFlow contrib module will not be included in TensorFlow 2.0.

For more information, please see:

- * <https://github.com/tensorflow/community/blob/master/rfcs/20180907-contrib-sunset.md>

- * <https://github.com/tensorflow/addons>

- * <https://github.com/tensorflow/io> (for I/O related ops)

If you depend on functionality not listed there, please file an issue.

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/contrib/predictor/saved_model_predictor.py:153: load (from tensorflow.python.saved_model.loader_impl) is deprecated and will be removed in a future version.

Instructions for updating:

This function will only be available through the v1 compatibility library as tf.compat.v1.saved_model.loader.load or tf.compat.v1.saved_model.load. There will be a new function for importing SavedModels in Tensorflow 2.0.


INFO:tensorflow:Restoring parameters from checkpoints/ckpt1635189825/variables/variables

In [19]:


```

for evaldata in evaltemps:
    prime_data = evaldata[OFFSET:OFFSET+PRIMELEN]
    results = prediction_run(predict_fn, prime_data, RUNLEN)
    utils_display.picture_this_6(evaldata, evaldates, prime_data, results, PRIME


```




RMSE on 109 predictions (shaded area): 8.488105536032977




RMSE on 109 predictions (shaded area): 4.599209103638949




RMSE on 109 predictions (shaded area): 10.158704121455557




RMSE on 109 predictions (shaded area): 11.773308889245957




RMSE on 109 predictions (shaded area): 4.778364026688139



RMSE on 109 predictions (shaded area): 5.6760404453718785



RMSE on 109 predictions (shaded area): 4.491668121229426



RMSE on 109 predictions (shaded area): 5.284793170023226

In [20]:

