

Create Keras Wide-and-Deep model

This notebook illustrates:

1. Creating a model using Keras. This requires TensorFlow 2.1

```
In [1]: # Ensure the right version of Tensorflow is installed.
!pip freeze | grep tensorflow==2.1
```

```
In [2]: # change these to try this notebook out
BUCKET = 'cloud-training-demos-ml'
PROJECT = 'cloud-training-demos'
REGION = 'us-centrall1'
```

```
In [3]: import os
os.environ['BUCKET'] = BUCKET
os.environ['PROJECT'] = PROJECT
os.environ['REGION'] = REGION
```

```
In [4]: %%bash
if ! gsutil ls | grep -q gs://${BUCKET}/; then
  gsutil mb -l ${REGION} gs://${BUCKET}
fi
```

Creating gs://cloud-training-demos-ml/...

ServiceException: 409 A Cloud Storage bucket named 'cloud-training-demos-ml' already exists. Try another name. Bucket names must be globally unique across all Google Cloud projects, including those outside of your organization.

```
-----
CalledProcessError                                Traceback (most recent call last)
<ipython-input-4-6b1d45d375e6> in <module>
----> 1 get_ipython().run_cell_magic('bash', '', 'if ! gsutil ls | grep -q gs://${BUCKET}/; then\n  gsutil mb -l ${REGION} gs://${BUCKET}\nfi\n')

/opt/conda/lib/python3.7/site-packages/IPython/core/interactiveshell.py in run_cell_magic(self, magic_name, line, cell)
    2401         with self.builtin_trap:
    2402             args = (magic_arg_s, cell)
--> 2403             result = fn(*args, **kwargs)
    2404         return result
    2405

/opt/conda/lib/python3.7/site-packages/IPython/core/magics/script.py in named_script_magic(line, cell)
    140         else:
    141             line = script
--> 142             return self.shebang(line, cell)
    143
    144         # write a basic docstring:

/opt/conda/lib/python3.7/site-packages/decorator.py in fun(*args, **kw)
    230         if not kwsyntax:
```

```

231         args, kw = fix(args, kw, sig)
--> 232         return caller(func, *(extras + args), **kw)
233     fun.__name__ = func.__name__
234     fun.__doc__ = func.__doc__

/opt/conda/lib/python3.7/site-packages/IPython/core/magic.py in <lambda>(f, *a,
**k)
185     # but it's overkill for just that one bit of state.
186     def magic_deco(arg):
--> 187         call = lambda f, *a, **k: f(*a, **k)
188
189         if callable(arg):

/opt/conda/lib/python3.7/site-packages/IPython/core/magics/script.py in shebang
(self, line, cell)
243         sys.stderr.flush()
244         if args.raise_error and p.returncode!=0:
--> 245             raise CalledProcessError(p.returncode, cell, output=out, std
err=err)
246
247     def _run_script(self, p, cell, to_close):

CalledProcessError: Command 'b'if ! gsutil ls | grep -q gs://{BUCKET}/; then\n
gsutil mb -l ${REGION} gs://{BUCKET}\nfi\n' returned non-zero exit status 1.

```

In []:

```

%%bash
ls *.csv

```

Create Keras model

First, write an input_fn to read the data.

In [7]:

```

import shutil
import numpy as np
import tensorflow as tf
print(tf.__version__)

```

2.3.3

In [8]:

```

# Determine CSV, label, and key columns
CSV_COLUMNS = 'weight_pounds,is_male,mother_age,plurality,gestation_weeks,key'.s
LABEL_COLUMN = 'weight_pounds'
KEY_COLUMN = 'key'

# Set default values for each CSV column. Treat is_male and plurality as strings
DEFAULTS = [[0.0], ['null'], [0.0], ['null'], [0.0], ['nokey']]

```

In [9]:

```

def features_and_labels(row_data):
    for unwanted_col in ['key']:
        row_data.pop(unwanted_col)
    label = row_data.pop(LABEL_COLUMN)
    return row_data, label # features, label

# load the training data
def load_dataset(pattern, batch_size=1, mode=tf.estimator.ModeKeys.EVAL):

```

```

dataset = (tf.data.experimental.make_csv_dataset(pattern, batch_size, CSV_COLUMNS,
          .map(features_and_labels) # features, label
        )
)
if mode == tf.estimator.ModeKeys.TRAIN:
    dataset = dataset.shuffle(1000).repeat()
dataset = dataset.prefetch(1) # take advantage of multi-threading; 1=AUTOTUNE
return dataset

```

Next, define the feature columns. `mother_age` and `gestation_weeks` should be numeric. The others (`is_male`, `plurality`) should be categorical.

In [11]:

```

## Build a Keras wide-and-deep model using its Functional API
def rmse(y_true, y_pred):
    return tf.sqrt(tf.reduce_mean(tf.square(y_pred - y_true)))

# Helper function to handle categorical columns
def categorical_fc(name, values):
    orig = tf.feature_column.categorical_column_with_vocabulary_list(name, value_list=values)
    wrapped = tf.feature_column.indicator_column(orig)
    return orig, wrapped

def build_wd_model(dnn_hidden_units = [64, 32], nembds = 3):
    # input layer
    deep_inputs = {
        colname : tf.keras.layers.Input(name=colname, shape=(), dtype='float32')
        for colname in ['mother_age', 'gestation_weeks']
    }
    wide_inputs = {
        colname : tf.keras.layers.Input(name=colname, shape=(), dtype='string')
        for colname in ['is_male', 'plurality']
    }
    inputs = {**wide_inputs, **deep_inputs}

    # feature columns from inputs
    deep_fc = {
        colname : tf.feature_column.numeric_column(colname)
        for colname in ['mother_age', 'gestation_weeks']
    }
    wide_fc = {}
    is_male, wide_fc['is_male'] = categorical_fc('is_male', ['True', 'False', 'Unknown'])
    plurality, wide_fc['plurality'] = categorical_fc('plurality',
        ['Single(1)', 'Twins(2)', 'Triplets(3)', 'Quadruplets(4)', 'Quintuplets(5)', 'Multiple(2+)'])

    # bucketize the float fields. This makes them wide
    age_buckets = tf.feature_column.bucketized_column(deep_fc['mother_age'],
        boundaries=np.arange(15, 45, 5))
    wide_fc['age_buckets'] = tf.feature_column.indicator_column(age_buckets)
    gestation_buckets = tf.feature_column.bucketized_column(deep_fc['gestation_weeks'],
        boundaries=np.arange(17, 47, 5))
    wide_fc['gestation_buckets'] = tf.feature_column.indicator_column(gestation_buckets)

    # cross all the wide columns. We have to do the crossing before we one-hot encode
    crossed = tf.feature_column.crossed_column(
        [is_male, plurality, age_buckets, gestation_buckets], hash_bucket_size=2**16)
    deep_fc['crossed_embeds'] = tf.feature_column.embedding_column(crossed, num_embeddings=nembds)

    # the constructor for DenseFeatures takes a list of numeric columns
    # The Functional API in Keras requires that you specify: LayerConstructor()

```

```

wide_inputs = tf.keras.layers.DenseFeatures(wide_fc.values(), name='wide_inp
deep_inputs = tf.keras.layers.DenseFeatures(deep_fc.values(), name='deep_inp

# hidden layers for the deep side
layers = [int(x) for x in dnn_hidden_units]
deep = deep_inputs
for layerno, numnodes in enumerate(layers):
    deep = tf.keras.layers.Dense(numnodes, activation='relu', name='dnn_{}'.
deep_out = deep

# linear model for the wide side
wide_out = tf.keras.layers.Dense(10, activation='relu', name='linear')(wide_

# concatenate the two sides
both = tf.keras.layers.concatenate([deep_out, wide_out], name='both')

# final output is a linear activation because this is regression
output = tf.keras.layers.Dense(1, activation='linear', name='weight')(both)
model = tf.keras.models.Model(inputs, output)
model.compile(optimizer='adam', loss='mse', metrics=[rmse, 'mse'])
return model

print("Here is our Wide-and-Deep architecture so far:\n")
model = build_wd_model()
print(model.summary())

```

Here is our Wide-and-Deep architecture so far:


Model: "functional_3"

Layer (type)	Output Shape	Param #	Connected to
gestation_weeks (InputLayer)	[(None,)]	0	
is_male (InputLayer)	[(None,)]	0	
mother_age (InputLayer)	[(None,)]	0	
plurality (InputLayer)	[(None,)]	0	
deep_inputs (DenseFeatures)	(None, 5)	60000	gestation_weeks [0][0] is_male[0][0] mother_age[0] [0] plurality[0][0]
dnn_1 (Dense)	(None, 64)	384	deep_inputs[0] [0]
wide_inputs (DenseFeatures)	(None, 71)	0	gestation_weeks [0][0] is_male[0][0]

			mother_age[0]
			plurality[0][0]
<hr/>			
dnn_2 (Dense)	(None, 32)	2080	dnn_1[0][0]
<hr/>			
linear (Dense) [0]	(None, 10)	720	wide_inputs[0]
<hr/>			
both (Concatenate)	(None, 42)	0	dnn_2[0][0] linear[0][0]
<hr/>			
weight (Dense)	(None, 1)	43	both[0][0]
<hr/>			
Total params: 63,227			
Trainable params: 63,227			
Non-trainable params: 0			
<hr/>			
None			

We can visualize the DNN using the Keras plot_model utility.

In [12]: `tf.keras.utils.plot_model(model, 'wd_model.png', show_shapes=False, rankdir='LR')`



Train and evaluate

In [13]:

```

TRAIN_BATCH_SIZE = 32
NUM_TRAIN_EXAMPLES = 10000 * 5 # training dataset repeats, so it will wrap around
NUM_EVALS = 5 # how many times to evaluate
NUM_EVAL_EXAMPLES = 10000 # enough to get a reasonable sample, but not so much t

trainds = load_dataset('train*', TRAIN_BATCH_SIZE, tf.estimator.ModeKeys.TRAIN)
evalds = load_dataset('eval*', 1000, tf.estimator.ModeKeys.EVAL).take(NUM_EVAL_E

steps_per_epoch = NUM_TRAIN_EXAMPLES // (TRAIN_BATCH_SIZE * NUM_EVALS)

history = model.fit(trainds,
                    validation_data=evalds,
                    epochs=NUM_EVALS,
                    steps_per_epoch=steps_per_epoch)

```

Epoch 1/5

312/312 [=====] - 4s 11ms/step - loss: 1.4109 - rmse: 1.1636 - mse: 1.4109 - val_loss: 1.2079 - val_rmse: 1.0988 - val_mse: 1.2079

```

Epoch 2/5
312/312 [=====] - 2s 8ms/step - loss: 1.2191 - rmse: 1.0933 - mse: 1.2191 - val_loss: 1.1369 - val_rmse: 1.0659 - val_mse: 1.1369
Epoch 3/5
312/312 [=====] - 3s 8ms/step - loss: 1.1486 - rmse: 1.0600 - mse: 1.1486 - val_loss: 1.1307 - val_rmse: 1.0631 - val_mse: 1.1307
Epoch 4/5
312/312 [=====] - 2s 7ms/step - loss: 1.1150 - rmse: 1.0439 - mse: 1.1150 - val_loss: 1.3552 - val_rmse: 1.1639 - val_mse: 1.3552
Epoch 5/5
312/312 [=====] - 2s 6ms/step - loss: 1.1350 - rmse: 1.0536 - mse: 1.1350 - val_loss: 1.1151 - val_rmse: 1.0557 - val_mse: 1.1151

```

Visualize loss curve

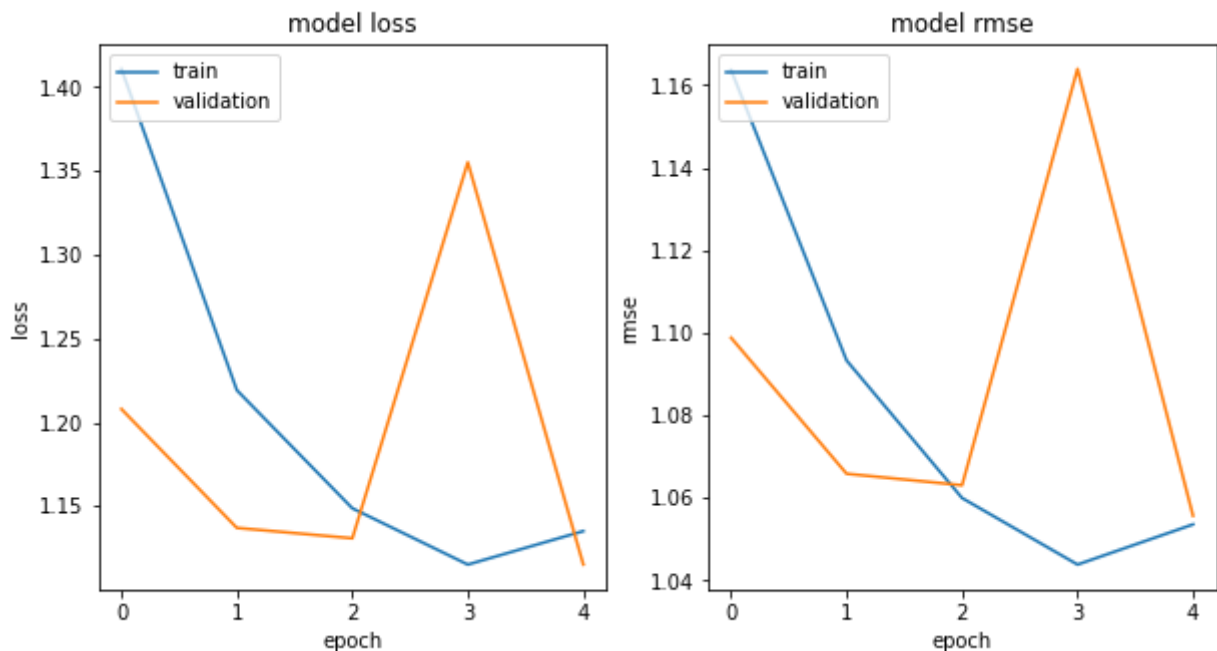
In [14]:

```

# plot
import matplotlib.pyplot as plt
nrows = 1
ncols = 2
fig = plt.figure(figsize=(10, 5))

for idx, key in enumerate(['loss', 'rmse']):
    ax = fig.add_subplot(nrows, ncols, idx+1)
    plt.plot(history.history[key])
    plt.plot(history.history['val_{}'.format(key)])
    plt.title('model {}'.format(key))
    plt.ylabel(key)
    plt.xlabel('epoch')
    plt.legend(['train', 'validation'], loc='upper left');

```



Save the model

In [15]:

```

import shutil, os, datetime
OUTPUT_DIR = 'babyweight_trained'
shutil.rmtree(OUTPUT_DIR, ignore_errors=True)
EXPORT_PATH = os.path.join(OUTPUT_DIR, datetime.datetime.now().strftime('%Y%m%d%

```

```
tf.saved_model.save(model, EXPORT_PATH) # with default serving function
print("Exported trained model to {}".format(EXPORT_PATH))
```

```
WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow/python/training/tracking/tracking.py:111: Model.state_updates (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.
Instructions for updating:
This property should not be used in TensorFlow 2.0, as updates are applied automatically.
WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow/python/training/tracking/tracking.py:111: Layer.updates (from tensorflow.python.keras.engine.base_layer) is deprecated and will be removed in a future version.
Instructions for updating:
This property should not be used in TensorFlow 2.0, as updates are applied automatically.
INFO:tensorflow:Assets written to: babyweight_trained/20210818183030/assets
Exported trained model to babyweight_trained/20210818183030
```

In [16]:

```
!ls $EXPORT_PATH
```

```
assets  saved_model.pb  variables
```

Copyright 2020 Google Inc. Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at <http://www.apache.org/licenses/LICENSE-2.0> Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License